

Knowing When to Abstain: Calibration and Selective Prediction for Machine-Learning Network Intrusion Detection

Anonymous (single author)
Affiliation withheld for review
author@example.org

Abstract—Machine-learning intrusion detectors are routinely deployed with a single hard decision and no notion of when their output should be trusted. We show, on the public NSL-KDD benchmark, that three standard detectors (logistic regression, a random forest, and gradient boosting) are markedly over-confident: expected calibration error reaches 0.2127, and the strongest model still attains an error rate of 0.2032 on a test set whose attack mix it was trained on. We then study two remedies. First, post-hoc Platt scaling fitted on in-distribution data fails to repair calibration once the test distribution shifts, leaving error essentially unchanged. Second, selective prediction—allowing the detector to abstain on its least-confident inputs—reduces selective risk from 0.2032 to 0.1136 at 80% coverage, but only for models whose confidence is informative; logistic regression’s confidence is nearly useless for ranking errors (area under the risk–coverage curve 0.2436 versus 0.0501). Finally, using a controlled family-holdout protocol we quantify behaviour on *unknown* attacks: detection of a held-out family collapses (e.g. DoS from 0.8696 to 0.6147), and abstention rejects only part of these unknowns—and almost none when the attack mimics benign traffic (R2L). On a second, modern corpus (CIC-IDS-2017) the within-distribution task is easy and well calibrated, yet under cross-day drift an unknown attack family is still missed and abstention rejects 0.9539 of it—so over-confidence is task-dependent while abstention generalises. We conclude that selective prediction is a cheap, model-agnostic safety layer for ML intrusion detection, with clearly characterised limits.

Index Terms—intrusion detection, selective prediction, calibration, uncertainty, open-set recognition, NSL-KDD

I. INTRODUCTION

A network intrusion detector that emits only “attack” or “benign” gives an operator no way to tell a confident verdict from a guess. In practice this matters twice over: alerts compete for scarce analyst attention, and genuinely novel attacks—absent from any training corpus—are exactly the cases a static classifier is least equipped to label. A detector that could say “I am not sure” on these inputs, and defer them to a human or a secondary control, would be more useful than one forced to commit.

This paper asks a narrow, practical question: do off-the-shelf machine-learning intrusion detectors produce confidence scores that are trustworthy enough to drive such an abstention policy, and what do they do when faced with attacks they have never seen? We answer empirically on NSL-KDD [1], a public benchmark whose published test split deliberately

contains attack types absent from training. Our contributions are:

- We measure the calibration of three standard detectors and show all are over-confident (Section IV-A), and that post-hoc Platt scaling [2] fitted on in-distribution data does *not* transfer to the shifted test set.
- We evaluate selective prediction [3], [4] and show abstention cuts selective risk substantially, but that its usefulness is model-dependent: tree ensembles rank their errors well, a linear model does not (Section IV-B).
- We introduce a controlled family-holdout protocol for the *unknown-attack* regime and show detection collapses on held-out families while abstention provides only partial cover—and almost none for attacks that resemble benign traffic (Section IV-C).

Throughout we follow the evaluation guidance of Arp *et al.* [5]: the calibrator is fit only on held-out training data, and no test information leaks into model selection. Unless stated otherwise, NSL-KDD test metrics are reported as the mean over $K=15$ seeds—resampling both the fit/calibration split and model randomness—with 95% confidence intervals (the canonical test partition is held fixed, so these intervals reflect train/model variance); equivalence is assessed with the *paired* per-seed difference and its CI rather than overlapping marginal intervals.

II. RELATED WORK

Calibration. Modern classifiers often output probabilities that do not match empirical accuracy [6], [7]. Expected calibration error (ECE) summarises this gap, and temperature/Platt scaling are common post-hoc fixes [2]. These analyses typically assume the evaluation data is drawn from the training distribution—an assumption that fails for intrusion detection.

Selective prediction. A classifier with a reject option [3] trades coverage for accuracy on what it does accept; the risk–coverage curve and its area [4] quantify the trade-off. Maximum softmax/posterior probability is a strong baseline confidence score [8].

ML for intrusion detection. The limits of learned detectors—especially on novel attacks and under distribution shift—were argued early [9] and remain a methodological hazard [5]. We do not propose a new detector; we ask whether existing ones know when they are wrong.

Unknown attacks and shift. The failure of closed-world classifiers on inputs unlike their training data motivates out-of-distribution and open-set detection, for which the maximum softmax probability is the standard baseline [8]. NSL-KDD operationalises this within a security benchmark: its test split deliberately contains attack types absent from training [1], so measuring behaviour on them needs no synthetic construction. Our family-holdout protocol complements that built-in novelty with a controlled version that makes the “unknown” condition causal. We deliberately reuse only off-the-shelf scores and standard metrics so that the conclusions concern the *evaluation* of detectors, not a new detector design.

III. METHODOLOGY

Data. NSL-KDD [1] provides 125973 training and 22544 test records with 41 features (three categorical) and a binary normal/attack label; each attack maps to one of four families (DoS, Probe, R2L, U2R). Its test split contains 17 attack *types* (3750 records) that never appear in training—a built-in open-set test. Numeric features are standardised and categorical features one-hot encoded; the encoder is fit on training data only.

Models. We use scikit-learn [10] logistic regression, a random forest, and histogram gradient boosting—deliberately standard detectors, since our subject is their confidence, not their architecture.

Calibration. For each model we report ECE (15 equal-width bins of the predicted class probability) before and after Platt scaling, the latter fitted on a 20% held-out split of the training set.

Selective prediction. The confidence of a prediction is its maximum class probability. Sorting the test set by confidence and accepting the most-confident fraction c yields the selective risk $r(c)$; we report the area under $r(c)$ (AURC) and $r(0.8)$.

Open-set protocol. To probe genuinely unknown attacks under control, we remove one attack family entirely from training (keeping benign traffic and the other families), retrain, and measure detection and confidence on that family’s test records. We also report, at a global threshold giving 80% coverage, the fraction of unknown-family records that abstention rejects.

IV. RESULTS

A. Detectors are over-confident

Table I reports accuracy and calibration. On the test set all three models are markedly over-confident, with ECE from ≈ 0.16 (random forest) to ≈ 0.22 (logistic regression). Crucially this is induced by distribution shift, not intrinsic: on a held-out *in-distribution* split the same models are nearly perfectly calibrated (random-forest ECE 0.0026, gradient-boosting 0.0005), and Platt scaling keeps them there (0.0005). It is only on the shifted test set—which introduces novel attacks—that ECE inflates (Fig. 3), and there post-hoc Platt scaling fit on source data cannot remove it—ECE stays essentially unchanged for gradient boosting and is slightly worse for the random forest. Calibrating to the source distribution does not certify probabilities once the deployed distribution moves (Fig. 2); this is the




Fig. 1. Risk–coverage curves on NSL-KDD. Tree ensembles rank errors well; the linear model does not.




Fig. 2. Reliability of the random forest, raw vs. Platt-scaled, on the shifted test set.

core reason we turn to abstention, which needs only a usable confidence *ranking*, not calibrated probabilities.

B. Abstention helps—if confidence is informative

Fig. 1 shows risk–coverage curves. For gradient boosting, abstaining on the least-confident 20% of inputs lowers selective risk to 0.1130 ± 0.0047 (AURC 0.0530 ± 0.0027); the random forest is similar (AURC 0.0604 ± 0.0013). Logistic regression is the cautionary case: its confidence barely ranks its errors (AURC 0.2259 ± 0.0115), so abstention buys little (0.1955 ± 0.0008). The lesson is that selective prediction is only as good as the underlying confidence signal, which here depends strongly on the model family.

C. Leave-one-family-out (LOFO) open-set evaluation

We formalize the unknown-attack test as a *leave-one-family-out* (LOFO) protocol: for each attack family F , remove all of F from training (keeping benign traffic and the other families), retrain, and measure detection and abstention on F ’s test records—a controlled, causal version of the genuine-unknown setting, which we release as reusable code. On NSL-KDD’s built-in novel attacks the random forest flags 0.7689 of *known*

TABLE I
NSL-KDD TEST METRICS: MEAN $\pm$ 95% CI OVER $K=15$ SEEDS (RESAMPLING THE FIT/CALIBRATION SPLIT AND MODEL RANDOMNESS). LOWER ECE AND AURC ARE BETTER.

Model	Acc.	ECE	AURC	$r(0.8)$
Logistic reg.	0.752 ± 0.001	0.2151 ± 0.0008	0.2259 ± 0.0115	0.1955 ± 0.0008
Random forest	0.769 ± 0.004	0.1630 ± 0.0036	0.0604 ± 0.0013	0.1328 ± 0.0026
Hist. grad. boost.	0.799 ± 0.004	0.1806 ± 0.0037	0.0530 ± 0.0027	0.1130 ± 0.0047




Fig. 3. Calibration error in-distribution vs. on the shifted test set. Models calibrated at training become badly miscalibrated under the test-time shift introduced by novel attacks.

TABLE II
LEAVE-ONE-FAMILY-OUT (LOFO) OPEN SET ON NSL-KDD: DETECTION WHEN A FAMILY IS SEEN IN TRAINING VS. HELD OUT (UNKNOWN), AND THE FRACTION OF UNKNOWN-FAMILY RECORDS ABSTENTION REJECTS AT 80% COVERAGE. U2R IS RARE ($n=52/67$), HENCE HIGH VARIANCE.

Family	det. (seen)	det. (unknown)	rejected by abstention
DoS	0.8696	0.6147	0.2918
Probe	0.7918	0.4374	0.3172
R2L	0.0586	0.0035	0.1737
U2R	0.1791	0.0299	0.6716

attacks but only 0.2336 of novel ones; LOFO (Table II, Fig. 4) confirms this causally: detection of DoS falls from 0.8696 to 0.6147 when DoS is unknown, and Probe from 0.7918 to 0.4374, with abstention rejecting 0.2918 and 0.3172 at 80% coverage. The rare U2R family (train $n=52$, test $n=67$ —a high-variance estimate) behaves like the anomalous families: abstention still flags 0.6716 of unknown U2R. R2L is the sharp exception—detected at only 0.0586 even when *seen*, 0.0035 when unknown, with abstention flagging just 0.1737. Because R2L traffic statistically resembles benign sessions the detector is confidently wrong, so the blind spot is specifically *benign-mimicking* attacks, not unknown attacks in general. The CIC-IDS-2017 cross-day experiment below is the same protocol on a second corpus, with the Web family held out.

D. Second dataset: CIC-IDS-2017 and cross-day drift

To test generality we repeat the study on CIC-IDS-2017, a modern flow-based corpus (78 features), using three days (DDoS, PortScan, and Web attacks). The picture is instruc-




Fig. 4. Detection rate per family when the family is seen in training vs. held out (unknown).

tively different (Table III). Within a random split the task is far easier than NSL-KDD—gradient boosting reaches accuracy 0.9998 at ECE 0.0002—so over-confidence is *not* a problem here: calibration quality is a property of how hard and how shifted the task is, not an intrinsic flaw of the models. The value of abstention, however, persists under drift. We evaluate it with a cross-day split—training and testing on different days rather than a random mix—which avoids the optimistic temporal bias a random split would introduce [11]. Training on Friday (DDoS, PortScan) and testing on Thursday—whose Web-attack family is unknown and whose benign traffic has drifted—detection of Web falls from 0.98 (seen) to 0.7812 (unknown); a confidence threshold at 80% coverage then rejects 0.9539 of those unknown-Web records, far more than the analogous NSL-KDD case because Web flows, unlike R2L, do not mimic benign traffic. The contrast across the two corpora is the substantive point: the calibration failure is task-dependent, while selective prediction is broadly useful—most when the unknown attack is statistically distinguishable from benign traffic.

E. Which uncertainty signal, and can we guarantee coverage?

Two questions sharpen the story: is maximum class probability (MSP) the right abstention signal, and can abstention carry a formal guarantee? For the first, we benchmark MSP against random-forest ensemble disagreement, Mahalanobis distance to class-conditional Gaussians, and k NN distance, on error ranking (AURC) and unknown-attack separation (Table IV). The appropriate statistic for a no-difference claim is the *paired* per-seed $\Delta\text{AURC} = \text{signal} - \text{MSP}$, and it

TABLE III
CIC-IDS-2017: REPLICATION (RANDOM SPLIT) AND CROSS-DAY DRIFT
(TRAIN FRIDAY, TEST THURSDAY).

Model	Acc.	ECE	AURC	$r(0.8)$
Logistic reg.	0.9893	0.0152	0.0014	0.0016
Random forest	0.9997	0.0007	0.0	0.0
Hist. grad. boost.	0.9998	0.0002	0.0	0.0

Cross-day: Web detection 0.98 $\rightarrow$ 0.7812 (unknown);
abstention rejects 0.9539 of unknown Web at 80% coverage.




Fig. 5. Risk-coverage on CIC-IDS-2017 (within-distribution split): the task is easy, so risk is low across coverages—unlike NSL-KDD (Fig. 1).

shows that no signal dominates. Ensemble disagreement is statistically indistinguishable from MSP (paired Δ AURC -0.0002 ± 0.0003 ; CI includes 0). Mahalanobis is marginally but significantly *better* at error ranking (-0.0022 ± 0.0014) yet significantly *worse* at unknown-attack detection (AUROC 0.807 ± 0.001 vs. MSP’s 0.881 ± 0.003); k NN is worse on both ($+0.0072 \pm 0.0030$). So contrary to the deep-learning out-of-distribution literature, where richer signals dominate under shift [12], [13], on tabular flow data the cheap max-softmax is a strong, well-rounded default—the leverage is in the abstention framework, not the estimator.

For the second, we wrap the detector in split conformal prediction [14], which gives distribution-free coverage under exchangeability. At target coverage 0.90, the prediction sets attain empirical coverage 0.930 ± 0.001 in-distribution—valid—but only 0.611 ± 0.002 under the train-test shift. This is the conformal analogue of the calibration failure of Section IV-A: the guarantee is exact under exchangeability and erodes precisely when the deployment distribution moves [15]. Restoring it needs shift-aware (weighted) conformal with estimable likelihood ratios—impractical for genuine zero-days, the honest caveat.

F. A benign-only novelty stage partially closes the blind spot

Confidence-based abstention barely flags R2L (Section IV-C) because such attacks lie on the benign feature manifold. The apt detector for that regime is a *benign-only* density model—a one-class novelty detector, distinct from




Fig. 6. Benign-only novelty detection (trained on benign records only): AUROC separating each attack family from benign. DoS/Probe are easy; R2L/U2R—the abstention blind spot—are only partially recoverable, R2L the hardest.

the two-class Mahalanobis used above for error ranking. We fit an Isolation Forest [16] and a one-class SVM [17] on benign records only. At a 10% benign false-positive budget the Isolation Forest recovers 0.537 of U2R and 0.24 of R2L (the one-class SVM 0.657 and 0.398)—against abstention’s ≈ 0.1737 for unknown R2L, a clear if partial gain. The AUROC ceilings (Fig. 6: R2L 0.824, U2R 0.892) show the limit is real but not absolute: R2L is hardest because its support most overlaps benign, so no benign-only detector fully recovers it, yet a novelty stage is a cheap, measured complement—abstention and novelty catch different failure modes (confident errors vs. benign-mimicking attacks).

V. DISCUSSION AND LIMITATIONS

Selective prediction is a cheap, model-agnostic safety layer: at the cost of deferring a minority of inputs it materially lowers the error among those acted upon, and it partially flags unknown attacks. But our results bound the claim. First, the value of abstention is inherited from the confidence signal, which is good for tree ensembles and poor for the linear model. Second, post-hoc calibration on source data does not survive distribution shift, so calibrated probabilities should not be trusted as deployment-time guarantees. Third, and most important operationally, abstention only catches unknowns that *look* anomalous; attacks that mimic benign traffic (R2L here) are missed confidently and silently.

Operational recipe. These results suggest a concrete deployment pattern. Pick a model whose confidence is informative—identified by a low AURC on validation data (here a tree ensemble, not the linear model); choose an operating coverage from the risk-coverage curve to meet a target precision on accepted alerts; route the abstained minority to a secondary control (a human analyst or an anomaly detector); and, because calibration does not survive shift, monitor the abstention rate over time as a cheap, label-free drift signal instead of trusting the raw probabilities—a lightweight complement to dedicated drift detectors [18], [19]. The reject option should be

TABLE IV

UNCERTAINTY SIGNALS ON NSL-KDD (MEAN $\pm$ 95% CI OVER $K=15$ SEEDS). AURC LOWER IS BETTER, UNKNOWN-AUROC HIGHER; Δ AURC IS THE PAIRED PER-SEED DIFFERENCE VS. MSP (NEGATIVE = BETTER; A CI SPANNING 0 MEANS NO DETECTABLE DIFFERENCE).

Signal	AURC	unknown-AUROC	Δ AURC vs. MSP
Max-softmax (MSP)	0.0604 $\pm$ 0.0013	0.881 $\pm$ 0.003	—
RF ensemble disagreement	0.0602 $\pm$ 0.0013	0.881 $\pm$ 0.003	-0.0002 $\pm$ 0.0003
Mahalanobis distance	0.0581 $\pm$ 0.0012	0.807 $\pm$ 0.001	-0.0022 $\pm$ 0.0014
k NN distance	0.0676 $\pm$ 0.0032	0.830 $\pm$ 0.006	+0.0072 $\pm$ 0.0030

paired with a complementary detector aimed at benign-looking intrusions (R2L/U2R), which abstention alone will not catch.

Limitations. NSL-KDD is dated and its features hand-crafted, so absolute numbers will differ on modern corpora; the mechanisms we isolate—over-confidence under shift, model-dependent confidence quality, and the benign-mimicry blind spot—are not specific to it. We study binary detection with a single abstention signal (maximum class probability); richer signals and the multi-class attack-typing setting are natural extensions. We evaluate CIC-IDS-2017 with a cross-day split; a *full* time-ordered evaluation (and the label-corrected CIC-IDS-2017 [20] and standardized NetFlow-v2 corpora for true cross-network transfer) remains future work.

VI. CONCLUSION

On a hard benchmark (NSL-KDD) off-the-shelf intrusion detectors are over-confident and post-hoc calibration does not transfer under shift; on an easier modern corpus (CIC-IDS-2017) they are well-calibrated, so the need for calibration is task-dependent. Across both, a simple reject option lowers selective risk and flags unknown attacks—strongly when those attacks are distinguishable from benign traffic, weakly when they mimic it. Selective prediction is worth adding to ML intrusion detection, provided its confidence signal and its blind spots are measured rather than assumed.

ACKNOWLEDGMENT

The authors used a large language model (Anthropic Claude) to assist with code scaffolding, editing, and literature search. All experiments, results, and claims were verified by the authors; the AI system is not an author. Generative-AI use is disclosed per IEEE policy.

REFERENCES

- [1] M. Tavallaei, E. Bagheri, W. Lu, and A. A. Ghorbani, “A detailed analysis of the KDD CUP 99 data set,” in *IEEE Symposium on Computational Intelligence for Security and Defense Applications (CISDA)*, 2009, pp. 1–6.
- [2] J. Platt, “Probabilistic outputs for support vector machines and comparisons to regularized likelihood methods,” in *Advances in Large Margin Classifiers*, A. J. Smola, P. Bartlett, B. Schölkopf, and D. Schuurmans, Eds. MIT Press, 1999, pp. 61–74.
- [3] R. El-Yaniv and Y. Wiener, “On the foundations of noise-free selective classification,” *Journal of Machine Learning Research*, vol. 11, pp. 1605–1641, 2010.
- [4] Y. Geifman and R. El-Yaniv, “Selective classification for deep neural networks,” in *Advances in Neural Information Processing Systems (NeurIPS)*, 2017, pp. 4878–4887.
- [5] D. Arp, E. Quiring, F. Pendlebury, A. Warnecke, F. Pierazzi, C. Wressnegger, L. Cavallaro, and K. Rieck, “Dos and don’ts of machine learning in computer security,” in *USENIX Security Symposium*, 2022, pp. 3971–3988.
- [6] C. Guo, G. Pleiss, Y. Sun, and K. Q. Weinberger, “On calibration of modern neural networks,” in *International Conference on Machine Learning (ICML)*, 2017, pp. 1321–1330.
- [7] A. Niculescu-Mizil and R. Caruana, “Predicting good probabilities with supervised learning,” in *International Conference on Machine Learning (ICML)*, 2005, pp. 625–632.
- [8] D. Hendrycks and K. Gimpel, “A baseline for detecting misclassified and out-of-distribution examples in neural networks,” in *International Conference on Learning Representations (ICLR)*, 2017.
- [9] R. Sommer and V. Paxson, “Outside the closed world: On using machine learning for network intrusion detection,” in *IEEE Symposium on Security and Privacy (S&P)*, 2010, pp. 305–316.
- [10] F. Pedregosa et al., “Scikit-learn: Machine learning in Python,” *Journal of Machine Learning Research*, vol. 12, pp. 2825–2830, 2011.
- [11] F. Pendlebury, F. Pierazzi, R. Jordaney, J. Kinder, and L. Cavallaro, “TESSERACT: Eliminating experimental bias in malware classification across space and time,” in *USENIX Security Symposium*, 2019, pp. 729–746.
- [12] Y. Ovadia, E. Fertig, J. Ren, Z. Nado, D. Sculley, S. Nowozin, J. V. Dillon, B. Lakshminarayanan, and J. Snoek, “Can you trust your model’s uncertainty? Evaluating predictive uncertainty under dataset shift,” in *Advances in Neural Information Processing Systems (NeurIPS)*, 2019.
- [13] K. Lee, K. Lee, H. Lee, and J. Shin, “A simple unified framework for detecting out-of-distribution samples and adversarial attacks,” in *Advances in Neural Information Processing Systems (NeurIPS)*, 2018, pp. 7167–7177.
- [14] A. N. Angelopoulos and S. Bates, “Conformal prediction: A gentle introduction,” *Foundations and Trends in Machine Learning*, vol. 16, no. 4, pp. 494–591, 2023.
- [15] R. J. Tibshirani, R. Foygel Barber, E. J. Candès, and A. Ramdas, “Conformal prediction under covariate shift,” in *Advances in Neural Information Processing Systems (NeurIPS)*, 2019, pp. 2526–2536.
- [16] F. T. Liu, K. M. Ting, and Z.-H. Zhou, “Isolation forest,” in *IEEE International Conference on Data Mining (ICDM)*, 2008, pp. 413–422.
- [17] B. Schölkopf, J. C. Platt, J. Shawe-Taylor, A. J. Smola, and R. C. Williamson, “Estimating the support of a high-dimensional distribution,” *Neural Computation*, vol. 13, no. 7, pp. 1443–1471, 2001.
- [18] R. Jordaney, K. Sharad, S. K. Dash, Z. Wang, D. Papini, I. Nouretdinov, and L. Cavallaro, “Transcend: Detecting concept drift in malware classification models,” in *USENIX Security Symposium*, 2017, pp. 625–642.
- [19] L. Yang, W. Guo, Q. Hao, A. Ciptadi, A. Ahmadzadeh, X. Xing, and G. Wang, “CADE: Detecting and explaining concept drift samples for security applications,” in *USENIX Security Symposium*, 2021, pp. 2327–2344.
- [20] G. Engelen, V. Rimmer, and W. Joosen, “Troubleshooting an intrusion detection dataset: The CICIDS2017 case study,” in *IEEE Security and Privacy Workshops (SPW)*, 2021, pp. 7–12.